

# Electronic Trading for Programmers

## Part I: Low-latency Execution



Mathias Gaunard

Quantitative Development Lead

16th of January 2023

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

# Outline

---

- Electronic Trading
  - Trading
  - Matching

- Low-Latency
  - Why things take time
  - Why being fast matters
  - Measuring

- Exchange connectivity
  - Feeds
  - Protocol considerations

- Trading system architecture
  - Constraints
  - Networking
  - State of the union

- Part II sneak peek

# What is trading?

---

Exchanging assets with another party

- base/quote, often quote is a stable currency, e.g BTC/USD
- buy/sell base (sell/buy quote)
- usually done either as an investment or as speculation

Different ecosystems

- equities (stocks, indices on stocks, ETFs)
- fixed income (government bonds, corporate bonds)
- commodities (oil, metals, grain)
- currencies
- cryptocurrencies

# How can I trade?

---

## Financial instruments

- Trade assets outright (spot)
- Obtain a loan and trade against that loan (spot with margin account)
- Enter a contract with trade obligations at term (futures, perp swaps, CFD)
- Enter a contract with trade optionality at term (options, warrants)
- Smart contracts (blockchain-based enforcement)
- Exotic contracts (sophisticated legally-binding agreements)

## Different products

- Listed on public exchanges
- Broker-dealer products
- Over-the-Counter only

# Why trade?

---

## Price move prediction (alpha)

- fundamental analysis of product (long term)
- events, news (medium term)
- market trends, statistics etc. (short term)

## Connecting people

- arbitrage buy/sell flow
- collecting fees
- arbitrage different marketplaces
- arbitrage derivative instruments on same assets

# Why electronically?

---

## Larger pool of participants

- link venues across the world
- connect retail and professionals
- more competition, better prices

## Transparency

- Records of all transactions
- Enforcement of due process
- MIFID compliance

## Automation

- Enables looking at small opportunities a human wouldn't consider
- Systematic algorithms to run strategies consistently
- Low-touch enables higher volume

# Who's trading?

---

## Investors, buy-side

- Pension funds, mutual funds
- Venture capitalists
- Hedge funds
- Proprietary trading firms
- Retail

## Trading services, sell-side

- Exchanges
- Market-makers
- Investment banks
- Brokers



# How to interact

---

## Direct, Over-the-Counter

- voice
- electronic, Request-for-Quote

## Through marketplace/exchange

- best participant selected (usually anonymous)
- small fees
- multiple platforms
  - continuous, "the screen"
  - auctions
  - multi-participant OTC-type platforms

## Execution on-behalf

- finds best way to enter large positions over longer time periods
- larger fees
- methodology pre-agreed and/or performance-tracked
  - Flow traders
  - Algo-driven, vwap/twap

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

# Vocabulary

---

- bid: buy order
- ask: sell order
- offer: ask
- side: whether it's buy or sell
- tick: increment of prices that are valid to place orders on
- prices  $p_1$  **better** than  $p_2$ :  $p_1 > p_2$  if bid,  $p_1 < p_2$  if ask
- **BBO**: best bid and ask
- bid-ask **spread**:  $best\_ask_{price} - best\_bid_{price}$
- bid and ask orders are **crossed**:  $bid_{price} \geq ask_{price}$
- **liquidity**: quantity addressable for trading, implied at good prices
- **touch**: liquidity close to the BBO

# Continuous matching

---

## Continuous

- Buyer/seller gets matched with sellers/buyers immediately if possible
  - If triggering match called **aggressor**, **taker** or **active** order
  - Removes matched liquidity, involved parties have traded
- Otherwise stays in order book and becomes **resting**, **maker** or **passive**
- Order book is always uncrossed
- Participants can amend/cancel their open orders

## Limit orders

- Instrument identifier and side
- Maximum quantity (number of lots)
- Worst price per lot

## Orders flags

- Immediate-or-Cancel
- Book-or-Cancel
- Icebergs


# Order book, initial state

## Order book

- Steady-state, all bids strictly less than asks
- Multiple orders per price level, arranged per insertion order (priority)

## Example scenario


- Tick of 0.50
- Spread of 1, e.g. two ticks



## Insertion example, join

Buy 20@10.50

- Join the queue on best bid
- Spread unaffected, two ticks wide




# Order book, back to initial



## Insertion example, improve


Buy 20@11

- Establish new price level
- Front of the queue
- Spread tightened to one tick






# Order book, back to initial




# Insertion example, take and improve

Buy 20@11.50

- Trade 10@11.50
- Sell order disappears
- Establish new price level at 11.50 with remaining quantity
- Front of queue
- Spread unaffected, but market "ticked up"




# Order book, back to initial



## Insertion example, take and widen

Buy 20@12

- Trade 10@11.50
- Sell order disappears
- Trade 10@12
- Buy order fully filled, not entering the book
- Sell order partially filled
- Spread widened, two ticks wide



# Order book, data structures

---

## Index orders by identifier

- Support modify and cancel
- Hash table

## Track priority of orders

- Linked-list of orders per level
- Ordered sequence of levels
  - Self-balancing binary tree
  - Circular buffer, dense tick representation

## Many updates per second

- Pre-allocate and pool
- Hybrid data structures
- Optimize for operations close to touch
- Optimize hashing for indexing method of exchange

# Outline

---

Electronic Trading

- Trading
- Matching

Low-Latency

- Why things take time
- Why being fast matters
- Measuring

Exchange connectivity

- Feeds
- Protocol considerations

Trading system architecture

- Constraints
- Networking
- State of the union

Part II sneak peek

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time**
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

# Asynchronous change requests

---

## Constant stream of updates

- Many participants
- Orders being added/modified/cancelled multiple times per millisecond
- Exchange has to distribute a lot of data, may be slow
- May even be throttled on its way to you

## Join the queue

- Observe market in a given state, want to affect a change
- Change request submitted and queued
- Exchange drains the queue of changes, eventually processes it
- Observe market with your change applied, tens of milliseconds might have passed



# Fairness

---

## Access to information

- Preferential access
- Does network provide information to all participants at the same time

## Order processing ordering

- Does everyone go through the same gateway
- Is there any reordering on the way to this gateway or internally

## Special rules make speed less of a concern

- Micro-auctions
- Asymmetric delays
- Pro-rata matching

# Determinism

---

## Deterministic

- Fastest always wins
- Clear, fair, efficient
- Technologically more challenging for exchange
- Leads to people building FPGAs and ASICs

## Non-deterministic

- Fastest only has an advantage, but semi-random
- Exchange can have dynamic behaviour based on load
- Leads to reverse-engineering and finding whatever can be gamified to improve the odds
- Mostly lots of tricks but no need for hardware

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters**
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

# Pick-offs

---

## Theoretical price-driven

- Only willing to buy for less/sell for more than some price model
- Market conditions change, theo changes, you cancel
- Aggressor is faster, bad trade

## Risk-driven

- Have lots of orders on many instruments or different venues
- Don't necessarily want them to be filled all at once
- One order is filled, cancel other ones
- Aggressor is faster, overtrade

## Achieve good priority

---

Bad priority  $\Rightarrow$  never get to trade

### Price improvement opportunities

- Event-driven and competitive
- Tied to price move predictions

Proactively place orders where price might move to

- More load and complexity
- Additional risk in the book

# Everyone wants to take the good stuff

---

## New order crosses the

- Spread tightened, new price attractive to your the
- Try and send an order to match against it
- Other participants might do the same, fastest takes it

## Event causes the change

- Something happens that affects your price model (usually other market)
- Now you think many existing orders are mispriced
- Try and send orders to match against all of them
- Other participants might do the same, may only get some of it

## Some metrics

---

### Fill rate

- How often do I get filled before cancelling
- Relative to various priority metrics

### Hitting rate

- How often do I get anything when I try to take liquidity
- Amount of credit/edge
- Latency of sending the message through

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek



# Time on the wire

---

## Trigger to order

- When was network packet triggering decision received
- When was network packet with order emitted out

## Streaming at 10Gbps

- MTU of 1500 is 1200ns
- End of Frame to End of Frame
- Start of Frame to Start of Frame
- Start of Trigger to Start of Frame

# Software times

---

## Not whole picture

- Delay between packet received and picked up by software
- Delay between sending packet and it being serialized out

## Low-latency measurements

- `rdtsc(p)`
- `rdpmc`

# Meaningful statistics

---

## Quantiles

- Min – how fast you can expect to get
- Median – how fast you are typically
- 90<sup>th</sup> percentile – how well are you protected
- 99<sup>th</sup> percentile – how bad can it get

## Determinism

- Removing tails is hard, but important
- Important events happen rarely
- Control flow divergence increases jitter

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

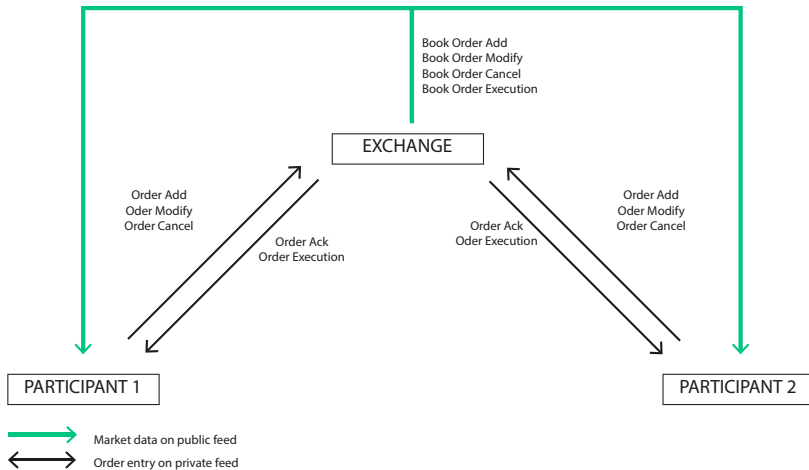
- Feeds**
- Protocol considerations

## Trading system architecture

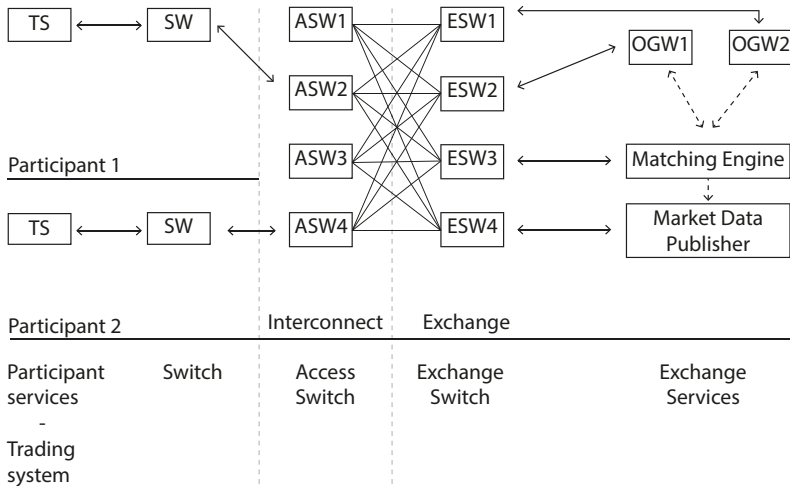
- Constraints
- Networking
- State of the union

## Part II sneak peek

# Participant connections



# Gateways and matching engine



# Unicast vs Multicast

---

## Unicast

- Send data for every single participant, bandwidth-hungry
- Goes through any router, including the open Internet
- No one gets it at the same time
- Can use TCP and have tailored per-participant data

## Multicast

- Send data once, switches fan-out, bandwidth-efficient
  - Requires ability to propagate subscribers through network
- Participants get data at the same time
  - Modulo network congestion and ethernet signal phase
- UDP-only



# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

# Order book fidelity

---

## Levels

- Level 1: BBO
- Level 2: aggregated quantity per price level
- Level 3: all individual orders

## Netting

- Unnetted
- Throttled
- Coalesced

# Recovery and reliability

---

## Sequence numbers

- Out of order
- Gaps
- Incremental updates

## Recovery

- Replay since beginning
- Snapshot

# Serialization formats

---

## Binary

- Flat, reinterpret\_cast-friendly formats, e.g. SBE
- Delta encoding, e.g. FAST

## Text

- FIX, key/value pairs
- JSON

## Fragmentation

- Nice exchanges avoid it
- Others whatever IP says goes

# Decimal numbers

---

## Problem

- Often working with non-integral prices and quantities
- 0.1 cannot be represented exactly with binary floating-point  
Find  $\{s, m, e\}$  such that  $v \simeq (-1)^s \times m \times 2^e$ , with  $1 \leq m < 2$
- Approximations cause all sorts of problems

## Decimal floating-point

- Find  $\{s, m, e\}$  such that  $v \simeq (-1)^s \times m \times 10^e$ , with  $1 \leq m < 10$
- IEEE754-2008, decimal32, decimal64, decimal128, backed by IBM and Intel
- Hardware support in POWER, software libraries
- Remains esoteric and slow

## Decimal numbers (2)

---

### Decimal fixed-point

- Fix the exponent and denormalize the mantissa
- Straightforward implementation, scaled integers
- Beware of operations that would change the scale
- No dynamic range/precision trade-off, beware of overflows

### Fixing approaches

- Compile-time exponent, part of type
- Runtime exponent, schema that applies to a dataset
  - no redundant storage with all values
  - e.g. time series

# Decimal numbers in practice

---

## Exchanges

- Some use decimal floating-point (rare)
- Most use decimal fixed-point with negative exponent
  - Either same exponent for everything on protocol
  - Or per-instrument exponent
- Some just use decimal text

## Recommendations

- Store data as integers, keep track of the relevant exponent they use
- Stick to scale-preserving operations when doing exact computations
- Switch in-and-out of double whenever doing non-exact numerical computations

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek



# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints**
- Networking
- State of the union

## Part II sneak peek

## Soft real-time constraints

---

By order of importance

- No system calls that block non-deterministically
- Real-time scheduling guarantees and affinity pinning
- Lock-free synchronization (i.e. no blocking system calls at all)
- No memory allocations
- No system calls at all
- Wait-free synchronization
- Limited code flow divergence (i.e. deterministic execution time)

# Hybrid trading systems

---

## Special-purpose

- Super fast cancel
  - Better cancel more often than needed
  - The simpler, the easier it is to make it fast
- Fast hitting
- Slower quoting, relaxed coding constraints

## Integrated

- More precise cancelling
- Enhanced capabilities for sophisticated strategies

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking**
- State of the union

## Part II sneak peek

# Ethernet, IP, TCP

---

## TCP – byte stream

- Handshake
- Ack window
- Nagle algorithm
- Retransmissions

## IP – packets

- Routing
- MTU and fragmentation

## Ethernet – frames

- Bandwidth
- Signal phase

# Low-latency networking

---

## Kernel-bypass

- Direct communication with Network Interface Adapter in userland
- DMA, write-combining memory, PCI-Express
- Disable interrupts, too slow to read means dropped data

## Userland TCP/IP

- Receive/send ethernet frames or frame fragments in application
- Re-implement all of TCP and IP without relying on the kernel
- Shortcuts for reliable networks

## Ready-made solutions

---

### Solarflare (now Xilinx)

- OpenOnload, BSD socket compatibility layer, high conformance
- EFVI, low-level API
- Onboard FPGA since Xilinx acquisition

### Exablaze (now Cisco)

- exanic low-level API, pre-loading capabilities
- Onboard FPGA as core of the system
- exasock, BSD socket compatibility layer, not-quite-conforming

### Others

- Mellanox (now Nvidia), more targeted at HPC and Infiniband
- Myricom, other HPC pioneer, notable for a Windows API, now defunct

# Standard approaches

---

## DPDK

- Linux foundation
- Large userland framework for kernel-bypass networking
- Supports traditional NICs (Intel, etc.)
- Userland TCP implementations

## io\_uring

- Linux kernel
- paradigm shift removing system calls
- new languages (e.g. Rust) building their networking around it



# Threading model

---

## Few threads

- No reliance on OS thread scheduling
- Cooperative scheduling intra-thread

## Some frameworks

- Asio, not the best fit
- Seastar, good principles

## Share-nothing

- Objects local to a given thread
- Communication via lock-free queues
- Seq-locks for state sampling
  - recently made compatible with C++ memory model

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union**

## Part II sneak peek

# State of the union

---

## Race to the bottom

- Normal software  $< 10ms$
- Software with real-time in mind  $< 100\mu s$  – sweet spot?
- Ultra low-latency software  $< 3\mu s$
- Normal FPGA solution  $< 500ns$
- Ultra low-latency FPGA solution  $< 50ns$
- ASIC  $< 30ns$
- Above and beyond  $< 10ns$

# Outline

---

## Electronic Trading

- Trading
- Matching

## Low-Latency

- Why things take time
- Why being fast matters
- Measuring

## Exchange connectivity

- Feeds
- Protocol considerations

## Trading system architecture

- Constraints
- Networking
- State of the union

## Part II sneak peek

## Part II sneak peek

---

### Pricing

- Pricing from trading activity
- Pricing derivatives

### Data analysis

- Time series
- Simulation

### Risk management

- Greeks
- Slippage

Questions?